



GraphRAG

By PK



Graph DB

A database built on **graph theory**

Data is stored as:

- **Nodes** → entities
- **Relationships (verbs)** → how entities connect
- **Properties** → attributes

Ex:

Nodes

Represent entities

Examples: Person, Movie, Product

Relationships

Connect two nodes

Examples: FOLLOWS, PURCHASED, LIKES

Properties

Key-value attributes

Examples: name: "Alice", age: 24

Nodes: Alice, Bob, Movie

Relationships:

Alice → FRIEND_OF → Bob

Bob → LIKES → Movie

Graph DB

Feature	Relational DB	Graph DB
Data model	Tables, Rows	Nodes, Relationships
Joins	Expensive	No joins; direct traversal
Schema	Rigid	Flexible
Deep query	Slow	Fast
Best for	Structured tabular	Connected data

- **Neo4j** is World's most popular **native graph database**
- Alternatives are **PuppyGraph, AWS Neptune, TigerGraph, OrientDB, and ArangoDB**
- Uses **Cypher Query Language**
- Visual graph exploration
- Offers AuraDB (cloud) for easy deployment



Neo4j Editions

Feature / Edition	Community	Desktop	Enterprise	Aura (Cloud)
Summary	Free basic Neo4j for small/local use	Enterprise features for free (dev only)	Full production-grade Neo4j	Fully managed cloud Neo4j
Cost	Free	Free	Paid	Free → Paid Aura Free (small) Aura Professional Aura Enterprise
Usage	Local + small production	Local development only	Production systems	Cloud-hosted apps
Clustering	✗ No	✗ No	✓ Yes	✓ Yes
Security (RBAC)	✗ No	✓ Yes (dev only)	✓ Full	✓ Full
Backups	Basic	Limited	✓ Hot backups	✓ Automatic backups
GUI	Neo4j Browser	Desktop App + Neo4j Browser	Neo4j Browser + tools	Web Console + Neo4j Browser
Cloud Managed	✗ No	✗ No	✗ No	✓ Yes
Best For	Learning, small apps	Developers, POCs	Large-scale production	Cloud apps, effortless deployment

Neo4j Architecture

- **Graph Engine**
 - The **Graph Engine** is the core part of Neo4j that stores and processes graph data.
- **Bolt Protocol** for fast access
 - Bolt is Neo4j's **binary client–server protocol**.
It is similar to JDBC but optimized specifically for graph queries, Provides fast, low-latency communication between your app and Neo4j
- **Write-Ahead Logging (WAL)**
 - A **Write-Ahead Log** is a file Neo4j writes **before** applying changes to the database and Only then does it update the graph store
- **Clustering:** Causal clustering
 - Neo4j's **distributed cluster architecture** that keeps data consistent across multiple machines.
Roles in the cluster
 - ◆ **Leader** → handles all writes
 - ◆ **Followers** → replicate data from leader & can handle reads
 - ◆ **Read Replicas** → scale out for massive read workloads
- **APOC(Awesome Procedures On Cypher):**
 - A huge library of **extensions** that provide advanced graph utilities and algorithms.
 - Data cleaning, Graph algorithms (similarity, pagerank, clustering), Import/export utilities, Text and date functions, ETL (Extract–Transform–Load) functions, Procedures to call external APIs. Graph transformations.



Cypher Query Language

Cypher Basics

1. Create a Node - Creates a Person node with a name property:

```
CREATE (p:Person {name: "Raj"})    or  MERGE (p:Person {name: "Raj"});
```

- Person is a Node label
- p is alias for Person node
- {name:"Raj"} is a property
- CREATE is used when you *always want a new node*.
- MERGE is used when you want a *unique node* (no duplicates).

2. Create a Relationship

Connects two existing people with a FRIEND_OF relationship:

```
MATCH (a:Person {name:"Alice"}), (b:Person {name:"Bob"})
CREATE (a)-[:FRIEND_OF]->(b)
```

- MATCH finds nodes
- CREATE adds the relationship
- Relationships have direction (→)



Cyper Query Language

2. Create a Relationship

Connects two existing people with a FRIEND_OF relationship:

```
CREATE (a:Person {name:"Alice"}), (b:Person {name:"Bob"})
CREATE (a)-[:FRIEND_OF]->(b)
```

- CREATE here create nodes and adds the relationship
- Relationships have direction (\rightarrow)

Pattern	Meaning	Direction?	Used in CREATE?
(-) [] - ()	Match relationship in ANY direction	No	✗ No
(-) [] -> ()	Outgoing relationship	Yes	✓ Yes
(-) <-[] - ()	Incoming relationship	Yes	✓ Yes



Cypher Query Language

3. Query

Find all friends of each person:

```
MATCH (p:Person)-[h:FRIEND_OF]->(friend)  
RETURN p,h, friend
```

- MATCH pattern-searches the graph
- Arrow means “connected by a FRIEND_OF relationship”
- Returns pairs of people and their friends

4. Query all movies an actor acted in

```
MATCH (p:Person {name: "Tom Hanks"})-[:ACTED_IN]->(m:Movie)  
RETURN m.title
```

- Finds the node for Tom Hanks
- Follows ACTED_IN relationships to movie nodes
- Returns titles of connected movies

```
MATCH p=()-[:DIRECTED]->()  
RETURN p LIMIT 25;
```

- Matches **any two nodes** connected by a **DIRECTED** relationship



Cypher Query Language

3. Update data

```
MATCH (p:Person {name:"Bob"})  
SET p.age = 26  
RETURN p;
```

4. Delete data

```
MATCH (p:Person {name:"Alice"})  
DETACH DELETE p;
```



Thank you